



ANÁLISE PREDITIVA DE POPULARIDADE DE NOTÍCIAS ONLINE

Arthur de Assis

Luciana Lima



ARTHUR DE ASSIS

Cientista de Dados na
A3Data

FORMAÇÃO

- 🎓 Sistemas de Informação
- 🎓 Mestrado em Ciência da Computação

INTERESSES

- 👍 Metaheurísticas
- 👍 Cozinhar
- 👍 Séries
- 👍 Machine Learning
- 👍 Twitter
- 👍 Lecionar

LUCIANA LIMA

Cientista de Dados na
A3Data

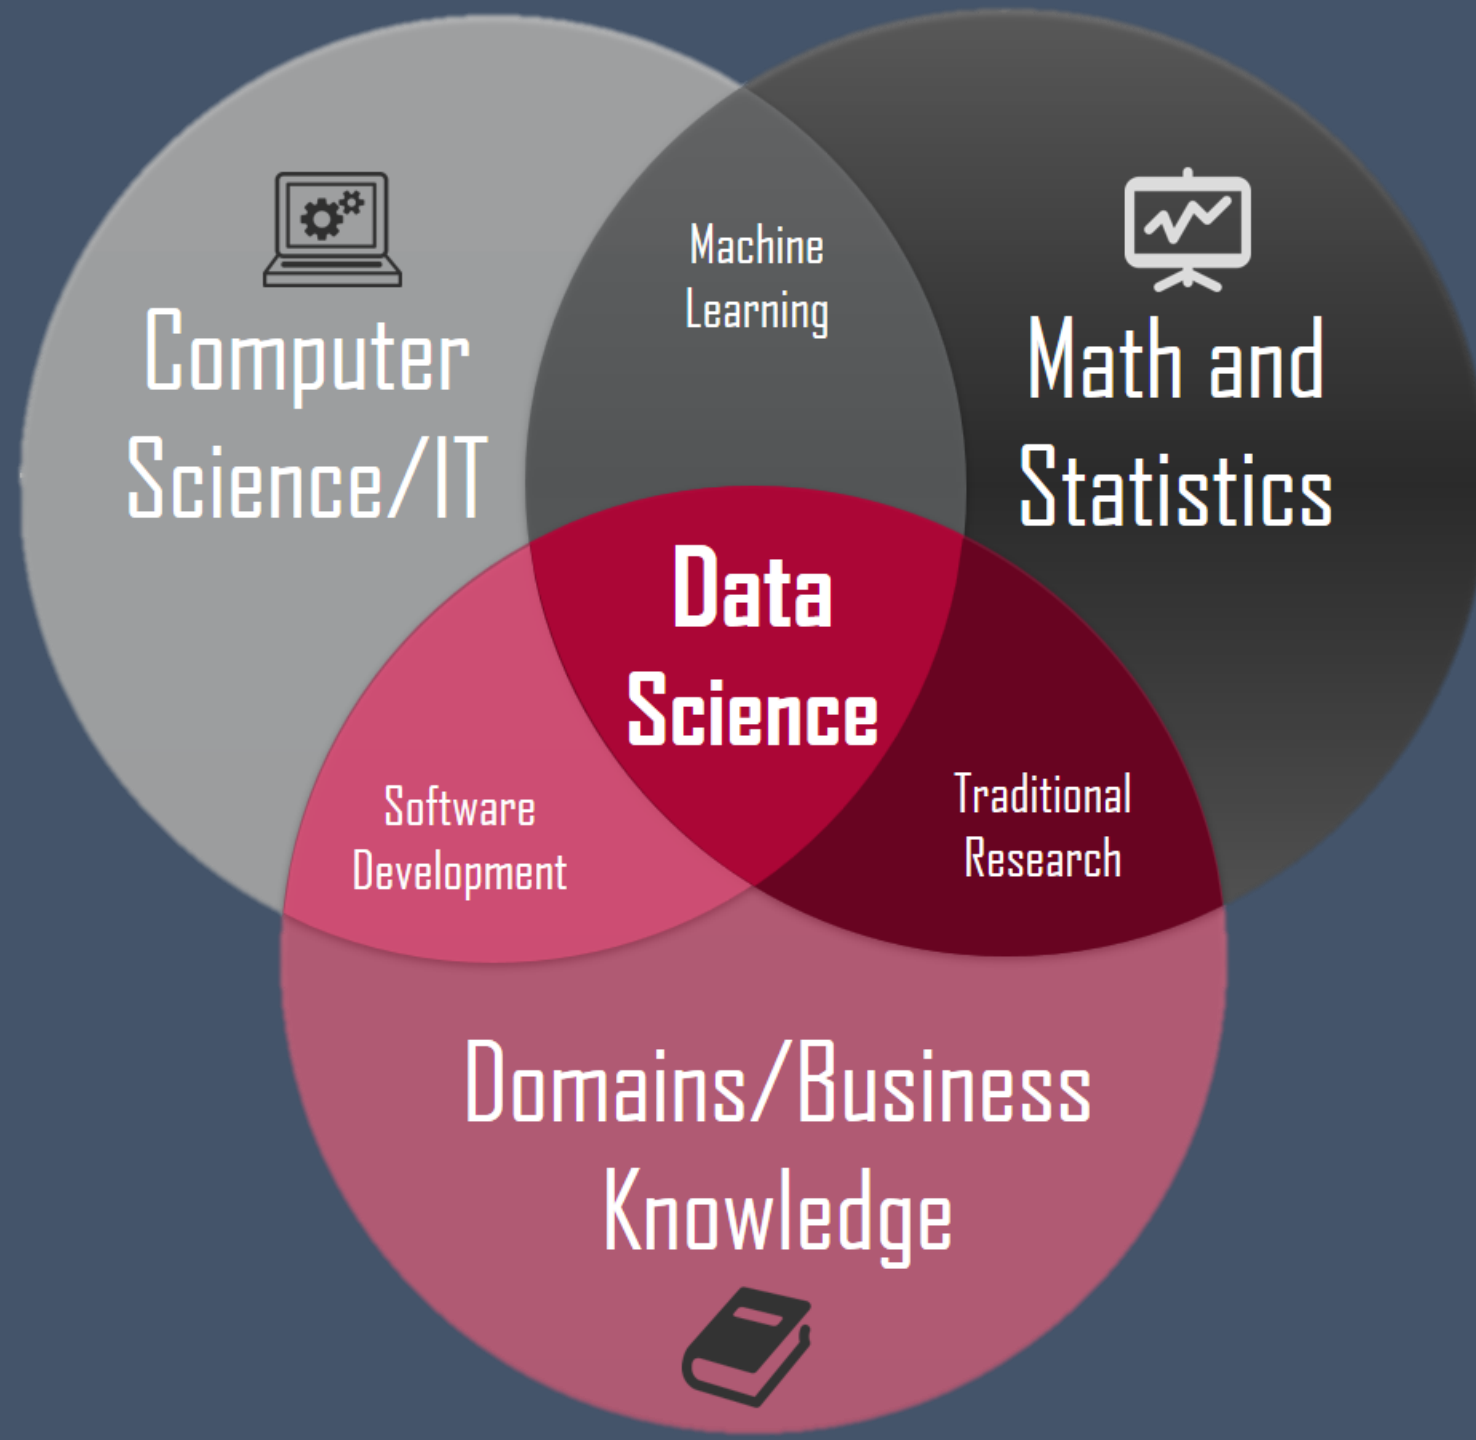
FORMAÇÃO

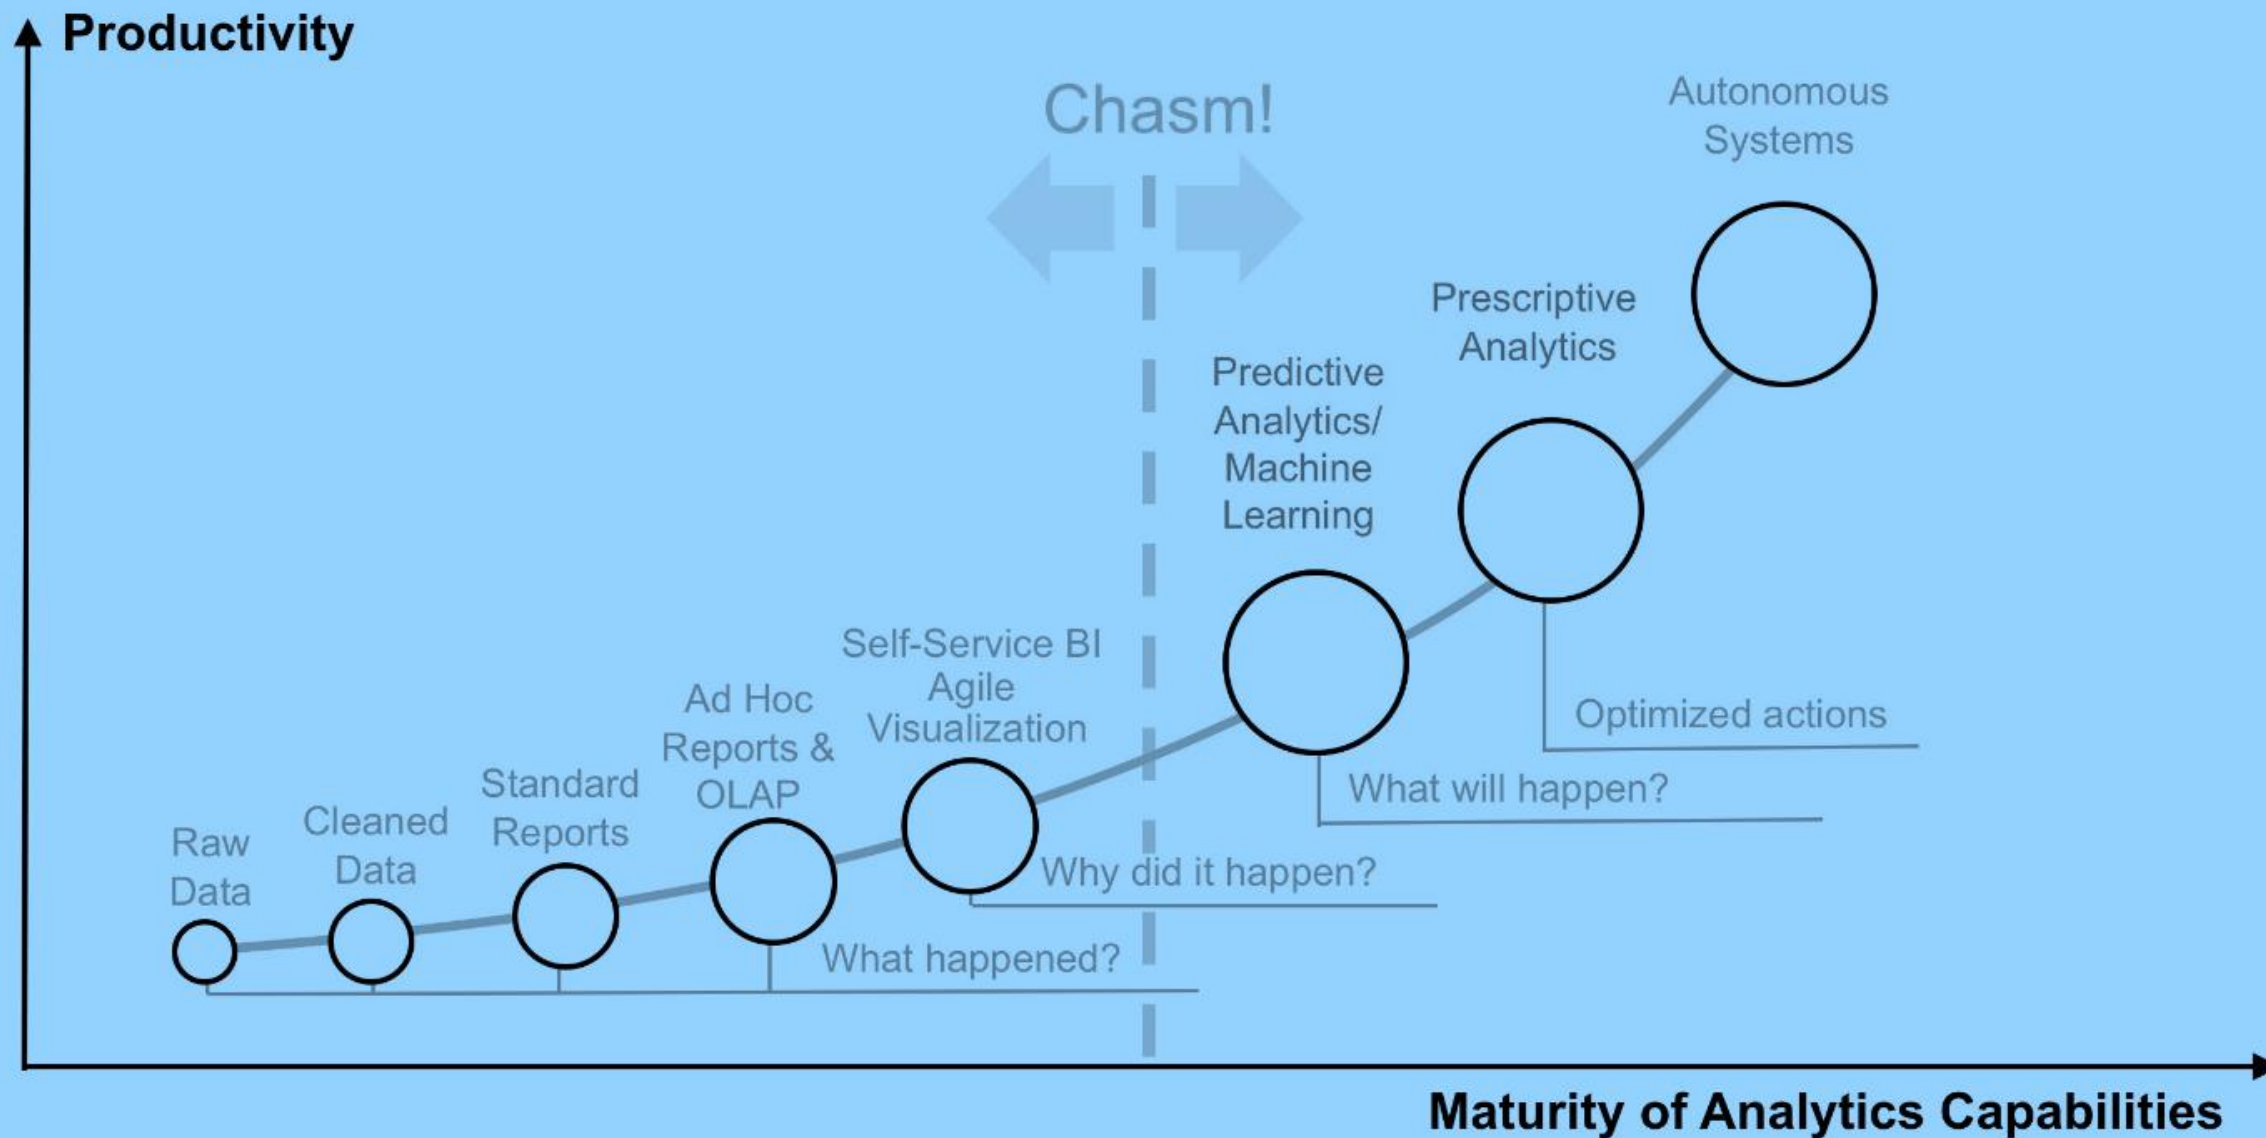
- Estatística 🎓
- Pós-graduação em Business Intelligence 🎓

INTERESSES

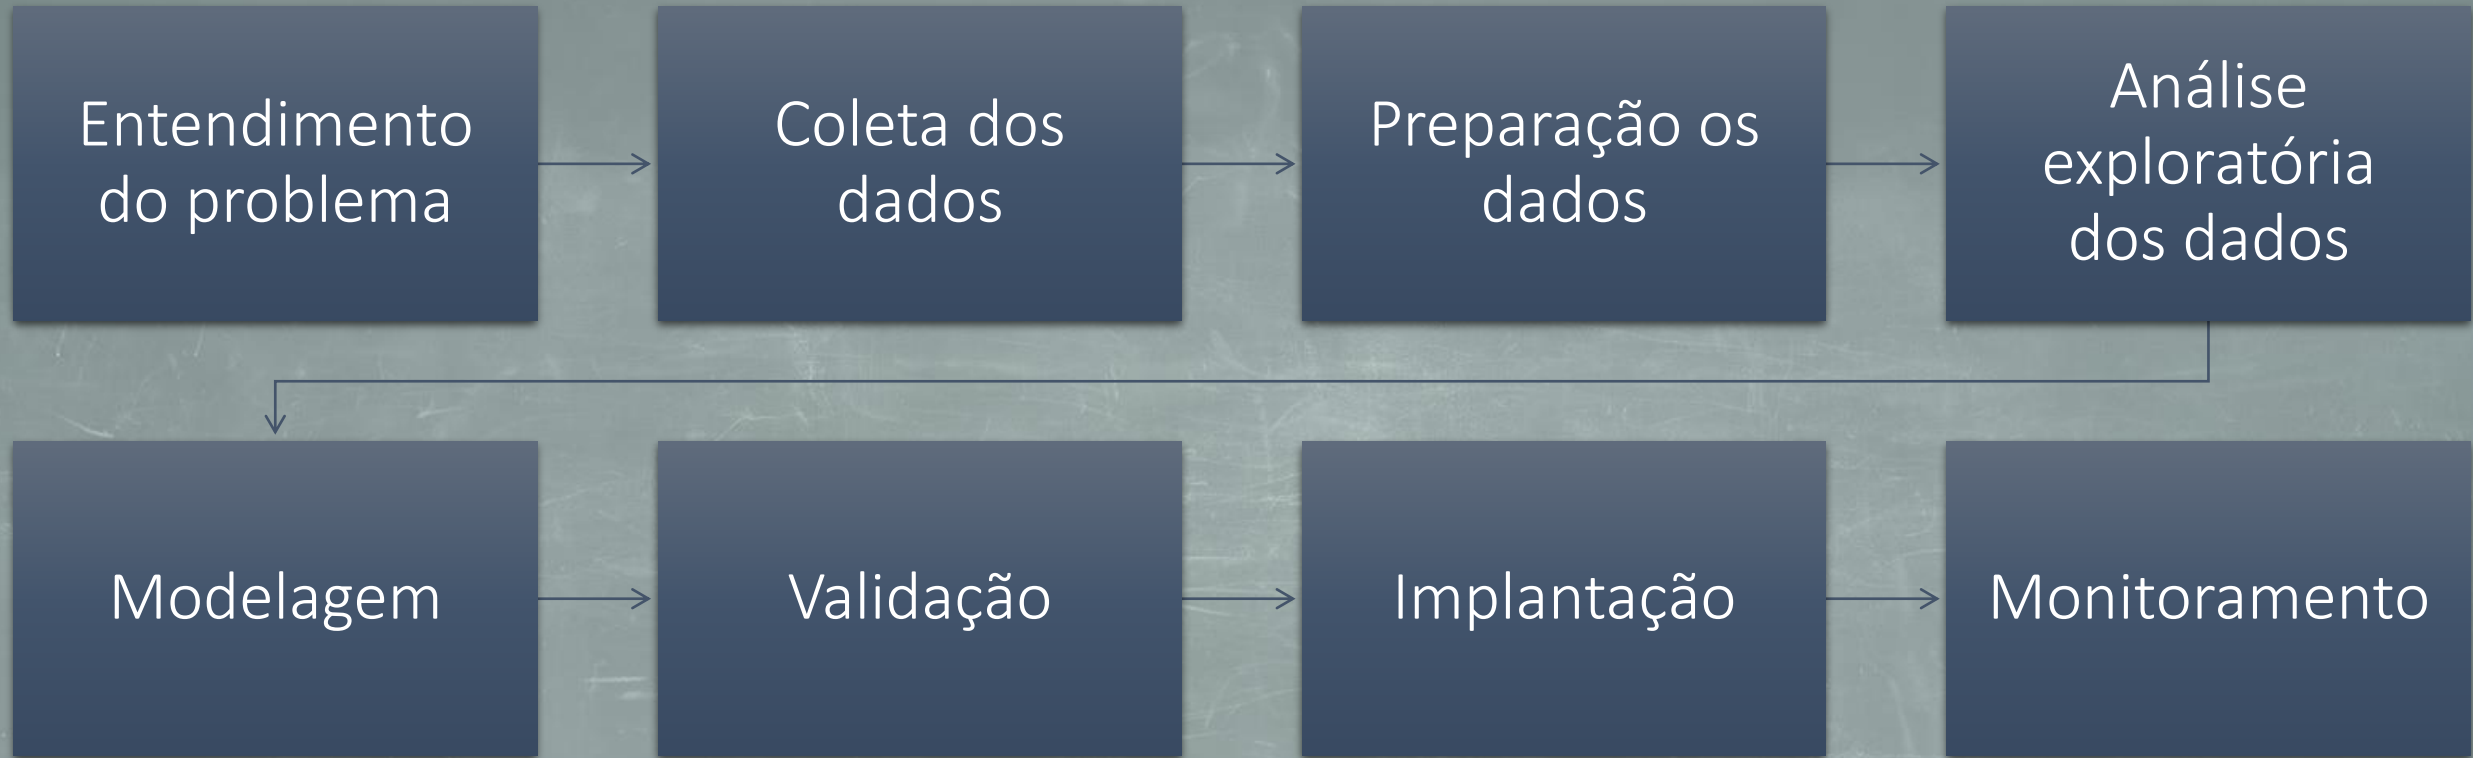
- Ciência de Dados 👍
- Futuro de Trabalho 👍
- Séries 👍
- Trekking 👍
- Viagem 👍
- Gatos 👍







8 PASSOS PARA UM PROJETO DE CIÊNCIA DE DADOS



What's New



Facebook's Privacy Disaster

3 days ago



Their house in London broke them up. So they bought a boat and brought up two kids on it.

3 days ago



YouTube Rewind 2018 is now the most disliked YouTube video of all time

3 days ago



Bees with sensor backpacks may help farmers monitor crops

3 days ago



MacBook Pro, Apple Watch, Microsoft Surface Pro 6, Instant Pot, KitchenAid, and more on sale for Dec. 13

3 days ago

What's Rising



TECH

Twitter releases 2018 Transparency Report including policy violation stats for the first time



What's Hot



CULTURE

Twitter thread explores questionable ways men, uh, decorate their bathrooms

Here are some genius ways to hang your toilet paper roll.



Days between the article publication and the dataset acquisition

Number of words in the title and content

Rate of unique and stopwords words in the content

Number of links, images and videos

Average length of the words in the content

Number of keywords in the metadata

Channel: Lifestyle, Entertainment, Business, Social Media, Tech, World?

Worst keyword (min, max, avg shares)

Best keyword (min, max, avg shares)

Avg. keyword (min, max, avg shares)

Referenced articles in Mashable (min, max, avg shares)

Day of week the article was published

Text subjectivity and sentiment polarity

Rate of positive and negative words in the content

Rate of positive and negative words among non-neutral tokens

Polarity of positive/negative words (min, max, avg)



Title subjectivity and sentiment polarity

Absolute subjectivity and sentiment polarity

VAMOS PRO
HANDS ON?



<https://colab.research.google.com/>

 Hello, Colaboratory 

File Edit View Insert Runtime Tools Help

+ CODE + TEXT

↑ CELL ↓ CELL

CONNECT ▾ EDITING

Table of contents Code snippets Files X

Getting Started

Highlighted Features

TensorFlow execution

GitHub

Visualization


Forms














Examples


Local runtime support

+ SECTION

EXAMPLES RECENT GOOGLE DRIVE GITHUB UPLOAD

Filter notebooks 

Title	First opened	Last opened	
 Hello, Colaboratory	Nov 23, 2018	0 minutes ago	
 online_news_popularity.ipynb	1 day ago	0 minutes ago	 
 Predict Shakespeare with Cloud TPUs and Keras	22 hours ago	22 hours ago	 
 External data: Drive, Sheets, and Cloud Storage	22 hours ago	22 hours ago	
 tensorflow - IGTI.ipynb	8 days ago	1 day ago	 

NEW PYTHON 3 NOTEBOOK  CANCEL

TensorFlow execution



```
[ ] !pip install -q tpot
    !pip install seaborn==0.9.0

    from sklearn.mixture import GaussianMixture
    from sklearn.cluster import KMeans
    from sklearn import metrics
    from scipy.spatial.distance import cdist
    import numpy as np
    import pandas as pd
    import matplotlib.pyplot as plt
    from sklearn import metrics
    from sklearn.metrics import pairwise_distances
    from sklearn.metrics.cluster import contingency_matrix
    from sklearn.decomposition import PCA
    from scipy.stats.mstats import zscore
    from pylab import savefig

    from sklearn.model_selection import train_test_split
    import seaborn as sns
    sns.set(style="white")
    sns.set(style="whitegrid", color_codes=True)

    from sklearn.metrics import r2_score

    from sklearn.linear_model import ElasticNet
    from sklearn.svm import SVR
    from sklearn.ensemble import RandomForestRegressor
    from sklearn.linear_model import LinearRegression
    from sklearn.model_selection import GridSearchCV

    from sklearn.neighbors import LocalOutlierFactor

    import requests
    import io

    from tpot import TPOTRegressor

    from sklearn.pipeline import make_pipeline, make_union
    from sklearn.preprocessing import StandardScaler
    from tpot.builtins import StackingEstimator
    from sklearn.linear_model import ElasticNetCV

    from google.colab import files # para fazer download dos arquivos gerados no Colab
```

```
[ ] class Preprocessor:

    def processor(self, df):

        self.cols_del = ['url', 'LDA_00', 'LDA_01', 'LDA_02', 'LDA_03', 'LDA_04']
        self.cols_norm = ['self_reference_max_shares', 'self_reference_min_shares',
                          'self_reference_avg_share', 'n_tokens_content', 'num_imgs',
                          'num_keywords', 'timedelta', 'n_tokens_title', 'num_hrefs',
                          'num_self_hrefs']
        df.columns = [col.strip() for col in df.columns]

        # Normaliza as columnas cols_norm
        self.stats = pd.DataFrame(columns = ["feature", "min", "max", "mean", "std"])
        for col in self.cols_norm:
            res = {"feature": col,
                  "min": df[col].min(),
                  "max": df[col].max(),
                  "mean": df[col].mean(),
                  "std": df[col].std()}
            self.stats = self.stats.append(res, ignore_index=True)
            df[col] = zscore(df[col])

        df['shares'] = df['shares'].apply(np.log10)

        df.drop(self.cols_del, axis=1, inplace=True)

    def reverter_zscore(self, feature, value):
        feat_values = pre.stats[pre.stats['feature'] == feature].iloc[0]
        return feat_values['std'] * value + feat_values['mean']
```



```
[ ] url="https://raw.githubusercontent.com/lgs coding/MachineLearning/master/data/OnlineNewsPopularity.csv"
df = pd.read_csv(url)
df.head()
```



	url	timedelta	n_tokens_title	n_tokens_content	n_unique_tokens	n_non_stop_words	n_non_stop_unique_tokens
0	http://mashable.com/2013/01/07/amazon-instant-...	731.0	12.0	219.0	0.663594	1.0	0.815385
1	http://mashable.com/2013/01/07/ap-samsung-spon...	731.0	9.0	255.0	0.604743	1.0	0.791946
2	http://mashable.com/2013/01/07/apple-40-billio...	731.0	9.0	211.0	0.575130	1.0	0.663866
3	http://mashable.com/2013/01/07/astronaut-notre...	731.0	9.0	531.0	0.503788	1.0	0.665635
4	http://mashable.com/2013/01/07/att-u-verse-apps/	731.0	13.0	1072.0	0.415646	1.0	0.540890

5 rows × 8 columns

```
[ ] df.shape
```

```
↳ (39644, 61)
```

```
[ ] df.apply(lambda a: a.isnull().sum(),axis=0)
```

```
↳
```

n_non_stop_unique_tokens	0
num_hrefs	0
num_self_hrefs	0
num_imgs	0
num_videos	0
average_token_length	0
num_keywords	0
data_channel_is_lifestyle	0
data_channel_is_entertainment	0
data_channel_is_bus	0
data_channel_is_socmed	0
data_channel_is_tech	0
data_channel_is_world	0
kw_min_min	0
kw_max_min	0
kw_avg_min	0
kw_min_max	0
kw_max_max	0
kw_avg_max	0
kw_min_avg	0
kw_max_avg	0
kw_avg_avg	0
self_reference_min_shares	0

```
[ ] pre = Preprocessor()  
pre.processor(df)  
df.head()
```



	timedelta	n_tokens_title	n_tokens_content	n_unique_tokens	n_non_stop_words	n_non_stop_unique_tokens
0	1.75788	0.757447	-0.695210	0.663594	1.0	0.815385
1	1.75788	-0.661657	-0.618794	0.604743	1.0	0.791946
2	1.75788	-0.661657	-0.712192	0.575130	1.0	0.663866
3	1.75788	-0.661657	-0.032933	0.503788	1.0	0.665635
4	1.75788	1.230482	1.115439	0.415646	1.0	0.540890

5 rows × 55 columns

▼ Função para plotar os outliers existentes na base

```
[ ] def plotar(df, grupo, x_label="x", y_label="y", titulo=None):  
    """  
    @param res_algoritmo : pd.DataFrame, deve ter as seguintes colunas (x, y, grupo)  
    """  
  
    # verifica quantos clusters tem  
    grupos = np.sort(df[grupo].unique())  
  
    plots = []  
    for g in grupos:  
        p = plt.scatter(df[df[grupo] == g].x, df[df[grupo] == g].y)  
        plots.append(p)  
  
    plt.legend(tuple(plots),  
                (tuple(["{g} {x}".format(g=grupo, x = str(c)) for c in grupos])),  
                loc=2, fontsize=8, bbox_to_anchor=(1.05, 1))  
    plt.xlabel(x_label)  
    plt.ylabel(y_label)  
    plt.title(titulo)  
    plt.show()
```

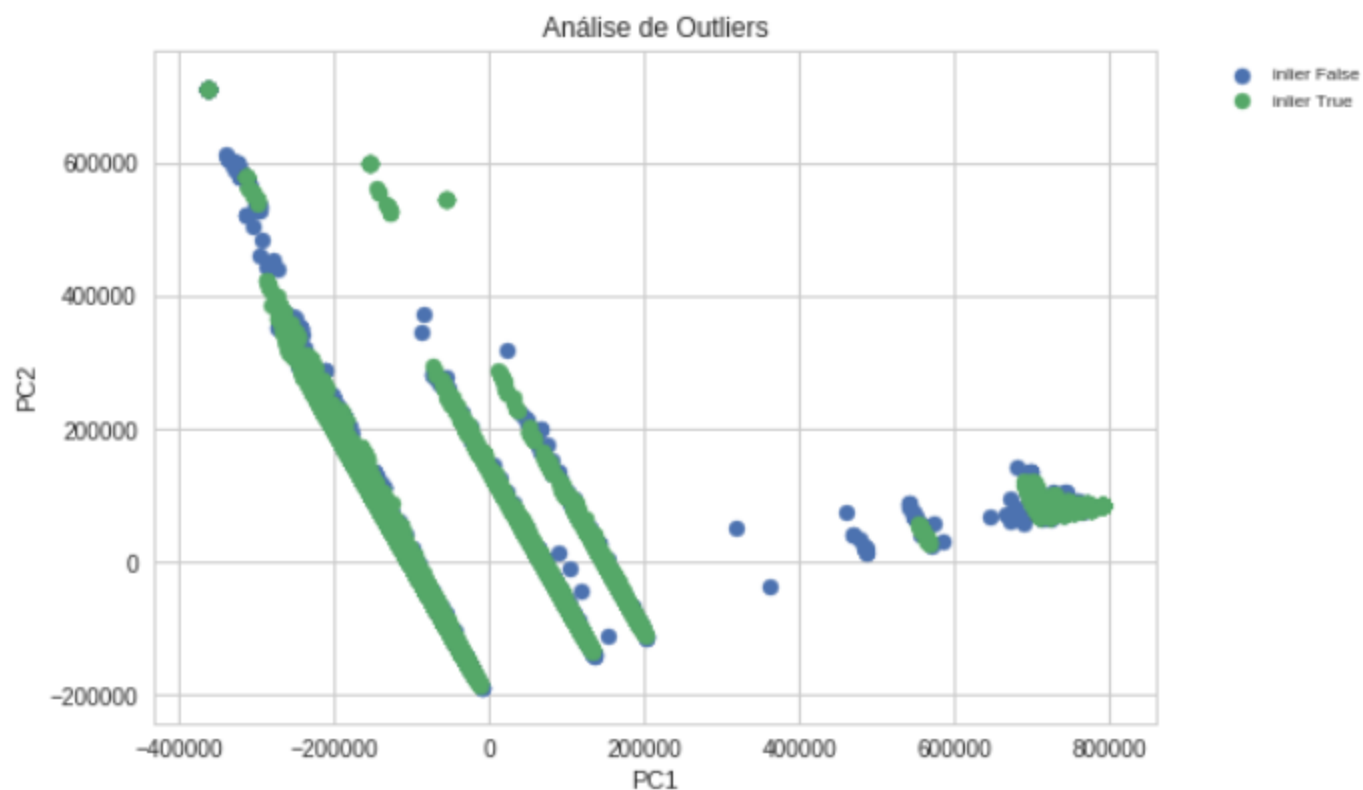
▼ Análise de Outliers

```
[ ] clf = LocalOutlierFactor(n_neighbors=20, contamination=0.1)
outlier = clf.fit_predict(df)
df['inlier'] = [True if x == 1 else False for x in outlier]

X = df.loc[:, df.columns != 'shares']
pca = PCA(n_components=2)
pca.fit(X)
pca_X = pca.transform(X)

df['x'] = pca_X[:,0]
df['y'] = pca_X[:,1]

plotar(df, "inlier", "PC1", "PC2", "Análise de Outliers")
```



▼ Filtrando os Outliers

```
[ ] df_sem_outlier = df[df['inlier'] == False]
df_sem_outlier.drop(['inlier', 'x', 'y'], axis = 1, inplace = True)
```

📄 /usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

▼ Separação das features e target

```
[ ] X = df_sem_outlier.loc[:, df_sem_outlier.columns != 'shares']
y = df_sem_outlier.loc[:, df_sem_outlier.columns == 'shares']
```


▼ Grid Search para os algoritmos que serão analisados:

- Random Forest Regressor
- ElasticNet
- SVM
- Regressão Linear

```
[ ] grid = {"elasticnet": {  
    'alpha': [0.5, 1.0, 1.5],  
    'l1_ratio': [0.25, 0.5, 0.75]  
},  
    "regressao_linear": {  
    'fit_intercept': [True, False]  
},  
    "svm": {  
    'kernel': ['rbf', 'linear', 'sigmoid'],  
    'C': [0.01, 0, 100]  
},  
    "random_forest": {  
    'max_features': ['auto', 'sqrt'],  
    'n_estimators': [200, 500]  
}  
}
```

```
[ ] algoritmos = {'elasticnet': ElasticNet(random_state=0, alpha = 1.5, l1_ratio = 0.75) ,  
    'svm': SVR(kernel = 'rbf', C = 100),  
    'regressao_linear': LinearRegression(fit_intercept = False),  
    'random_forest': RandomForestRegressor(random_state=0, max_features = 'sqrt', n_estimators=500)}
```

▼ Split da base em treino (70%) e teste (30%)

```
[ ] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)
```



```
resultados = pd.DataFrame(columns=['algoritmo', 'parametros', 'score_treino',  
                                   'score_teste', "evs", "mae", "mse", "msle"])  
  
predicts = y_test.reset_index()  
  
for nome_alg in algoritmos:  
    alg = algoritmos[nome_alg]  
    alg.fit(X_train, y_train)  
    y_pred = alg.predict(X_test)  
    predicts['pred_' + nome_alg] = y_pred  
    res = {"algoritmo": nome_alg,  
          "parametros": None,  
          "score_treino": alg.score(X_train, y_train),  
          "score_teste": alg.score(X_test, y_test),  
          "evs": metrics.explained_variance_score(y_test, y_pred),  
          "mae": metrics.mean_absolute_error(y_test, y_pred),  
          "mse": metrics.mean_squared_error(y_test, y_pred),  
          "msle": metrics.mean_squared_log_error(y_test, y_pred)}#None}  
  
resultados = resultados.append(res, ignore_index=True)
```

```
[ ] resultados
```



	algoritmo	parametros	score_treino	score_teste	evs	mae	mse	msle
0	elasticnet	None	0.061922	0.083592	0.083642	0.342678	0.207691	0.010581
1	svm	None	0.954667	0.013183	0.015164	0.360590	0.223648	0.011328
2	regressao_linear	None	0.125333	0.092393	0.092462	0.339157	0.205696	0.010443
3	random_forest	None	0.880729	0.134511	0.135064	0.332257	0.196151	0.009912

▼ Utilização do autoML para construção do melhor modelo

```
[ ] tpot = TPOTRegressor(generations=5, population_size=20, verbosity=2)
    tpot.fit(X_train, y_train)
    print(tpot.score(X_test, y_test))
    tpot.export('tpot_shares_pipeline.py')
    files.download('tpot_shares_pipeline.py')
```

```
[ ] pipeline = make_pipeline(
    StandardScaler(),
    StackingEstimator(estimator=ElasticNetCV(l1_ratio=0.8500000000000001, tol=0.0001)),
    RandomForestRegressor(bootstrap=True, max_features=0.8, min_samples_leaf=14, min_samples_split=13, n_estimators=100)
)
```

▼ Aplicação do melhor modelo gerador pelo AutoML

```
[ ] pipeline.fit(X_train, y_train)
    results = pipeline.predict(X_test)

    res = {"algoritmo": 'stacking',
          "parametros": None,
          "score_treino": pipeline.score(X_train, y_train),
          "score_teste": pipeline.score(X_test, results),
          "evs": metrics.explained_variance_score(y_test, results),
          "mae": metrics.mean_absolute_error(y_test, results),
          "mse": metrics.mean_squared_error(y_test, results),
          "msle": metrics.mean_squared_log_error(y_test, results)}

    resultados = resultados.append(res, ignore_index=True)
```



```
[ ] resultados
```

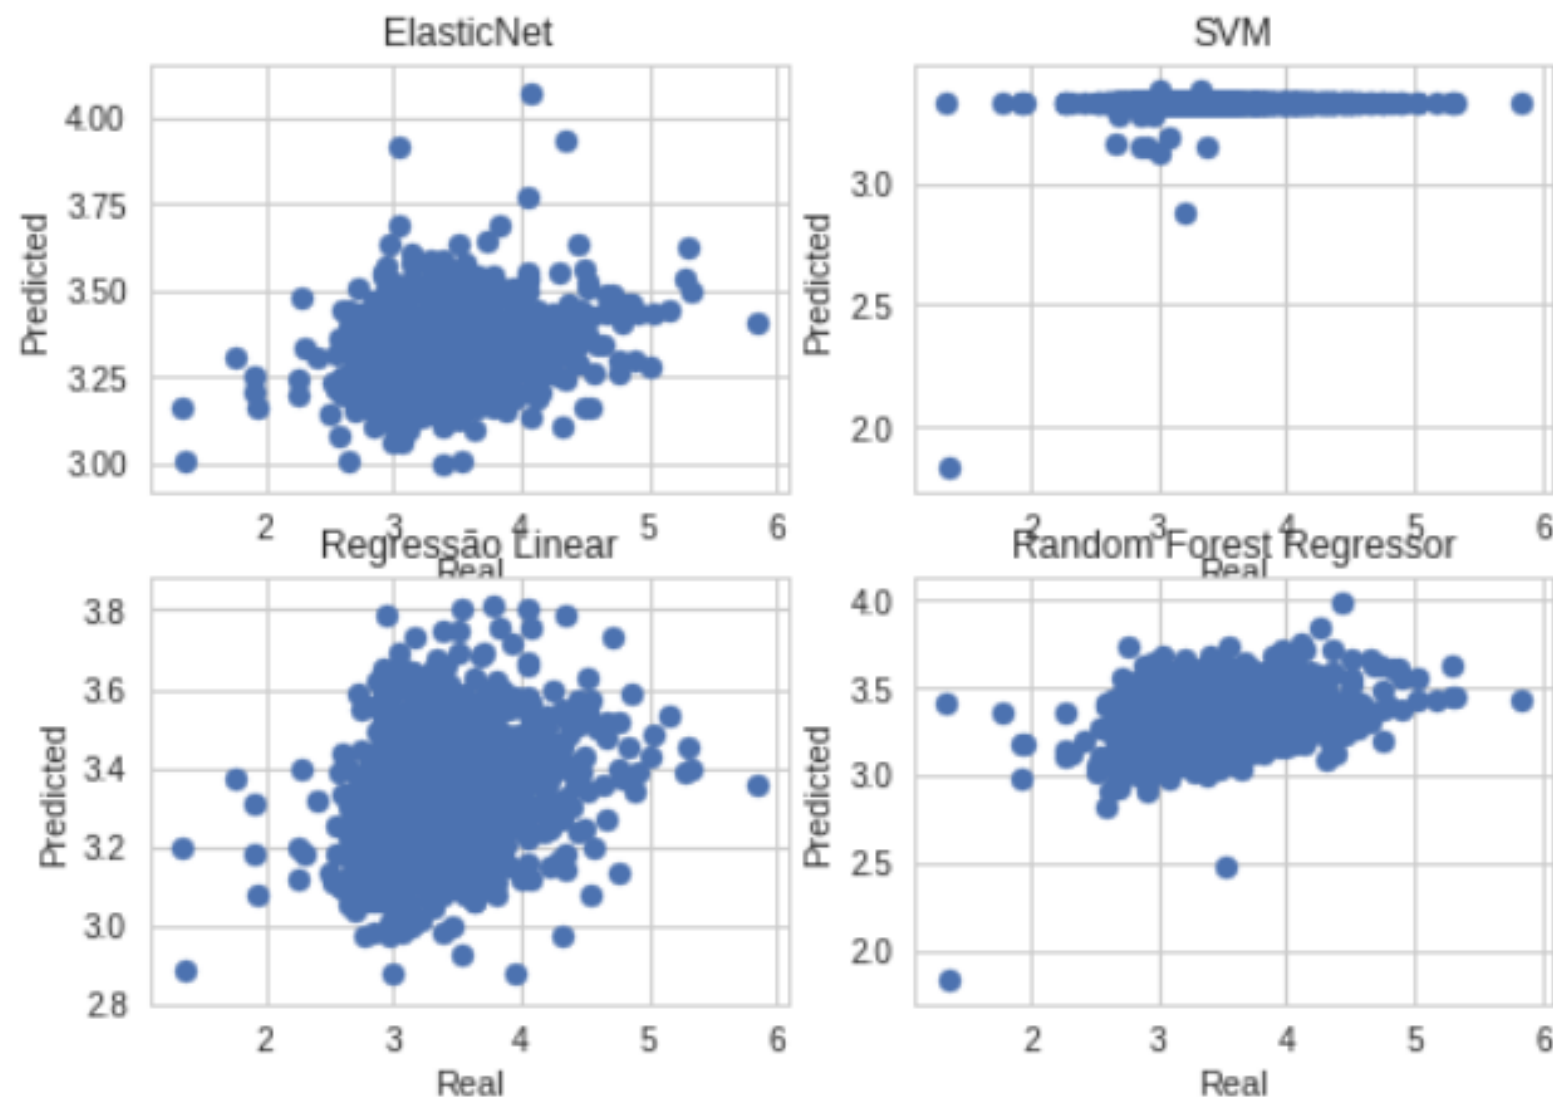


	algoritmo	parametros	score_treino	score_teste	evs	mae	mse	msle
0	elasticnet	None	0.061922	0.083592	0.083642	0.342678	0.207691	0.010581
1	svm	None	0.954667	0.013183	0.015164	0.360590	0.223648	0.011328
2	regressao_linear	None	0.125333	0.092393	0.092462	0.339157	0.205696	0.010443
3	random_forest	None	0.880729	0.134511	0.135064	0.332257	0.196151	0.009912
4	stacking	None	0.437629	1.000000	0.139023	0.329617	0.195148	0.009879
5	stacking	None	0.431962	1.000000	0.137278	0.331082	0.195557	0.009899



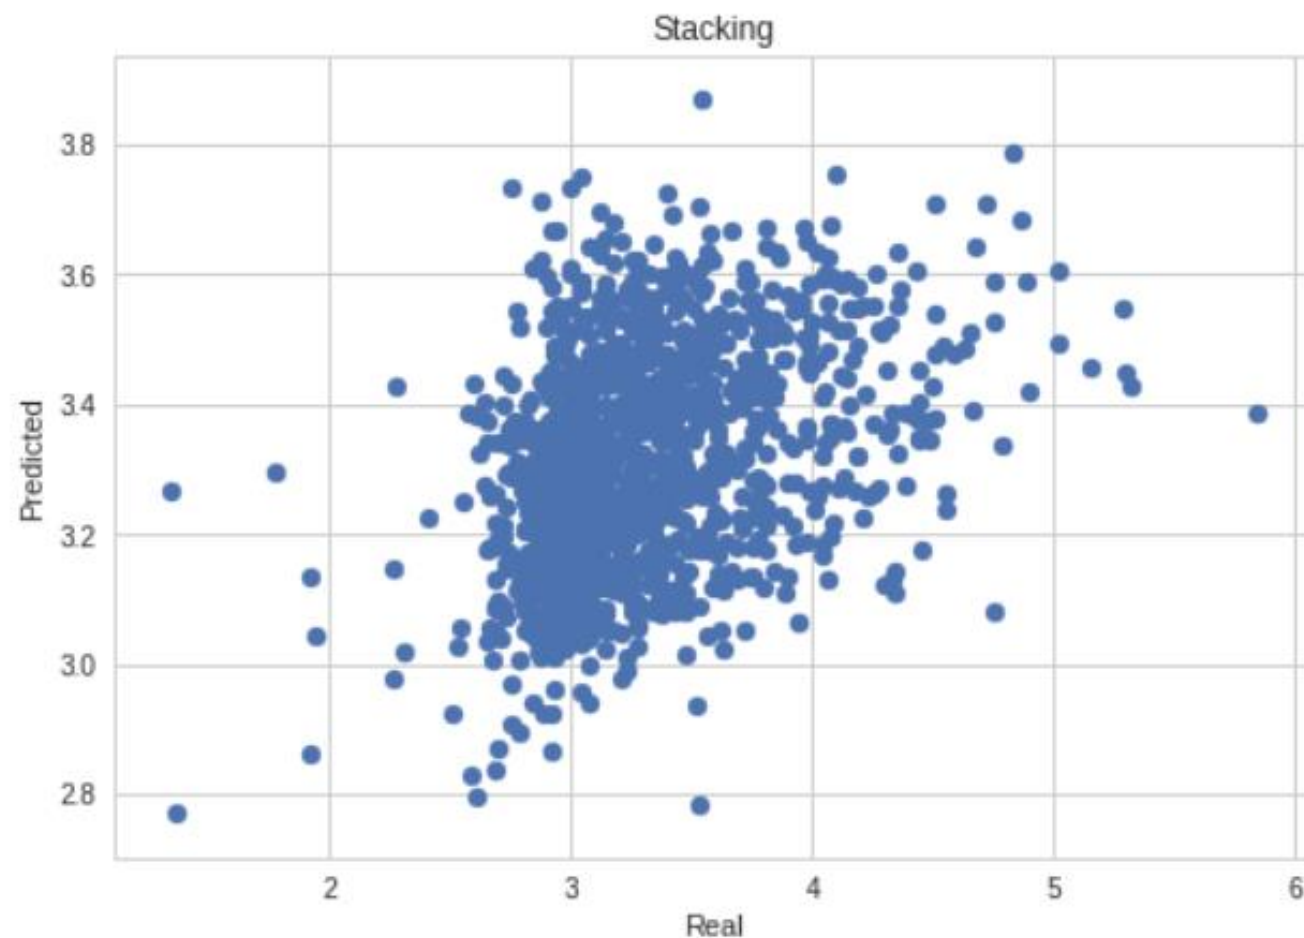
```
#@title
plt.subplot(2,2, 1)
en = plt.scatter(x="shares", y="pred_elasticnet", data=predicts)
plt.xlabel("Real")
plt.ylabel("Predicted")
plt.title("ElasticNet")
plt.subplot(2,2, 2)
sv = plt.scatter(x="shares", y="pred_svm", data=predicts)
plt.xlabel("Real")
plt.ylabel("Predicted")
plt.title("SVM")
plt.subplot(2,2, 3)
rl = plt.scatter(x="shares", y="pred_regressao_linear", data=predicts)
plt.xlabel("Real")
plt.ylabel("Predicted")
plt.title("Regressão Linear")
plt.subplot(2,2, 4)
rf = plt.scatter(x="shares", y="pred_random_forest", data=predicts)
plt.xlabel("Real")
plt.ylabel("Predicted")
plt.title("Random Forest Regressor")
```

↳ `Text(0.5,1,'Random Forest Regressor')`



```
[ ] predicts['pred_stacking'] = results
plt.scatter(x="shares", y="pred_stacking", data=predicts)
plt.xlabel("Real")
plt.ylabel("Predicted")
plt.title("Stacking").
```

➞ Text(0.5,1,'Stacking')



CONSIDERAÇÕES
FINAIS

OBRIGADO(A)!

arthur.assis@a3data.com.br
luciana.lima@a3data.com.br